

▼ Author : Braham Parkash

## Data Science & Business Analytics Internship

### GRIP - The Spark Foundation

**Task 2 : From the given 'Iris' dataset, predict the optimum number of clusters and represent it visually.**

Machine learning can be of four types- Supervised, Unsupervised, Semi-supervised and Reinforcement learning. Unlike Supervised learning, Unsupervised learning has features measured for n observations but has no response, i.e., we have unlabeled training data. In this situation we cannot predict a response (as there is no response variable) and therefore we are left with the task of finding interesting patterns in the data. Some of the important unsupervised learning algorithms are - Clustering, Principal Component Analysis(PCA) and Association Rule learning. In this kernel we will explore Clustering and PCA

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
data=pd.read_csv('IRIS.csv')
```

```
data.head()
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
```

```
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
data.describe()
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
data.duplicated().sum()
```

```
3
```

```
data.drop_duplicates(inplace=True)
```

```
data.duplicated().sum()
```

```
0
```

```
print('Data has {} number of rows and {} columns'.format(data.shape[0],data.shape[1]))
```

```
Data has 147 number of rows and 5 columns
```

```
print('The columns set in our dataset :',data.columns)
```

```
The columns set in our dataset : Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_wi  
'species'],  
dtype='object')
```

```
#Frequency distribution of species
```

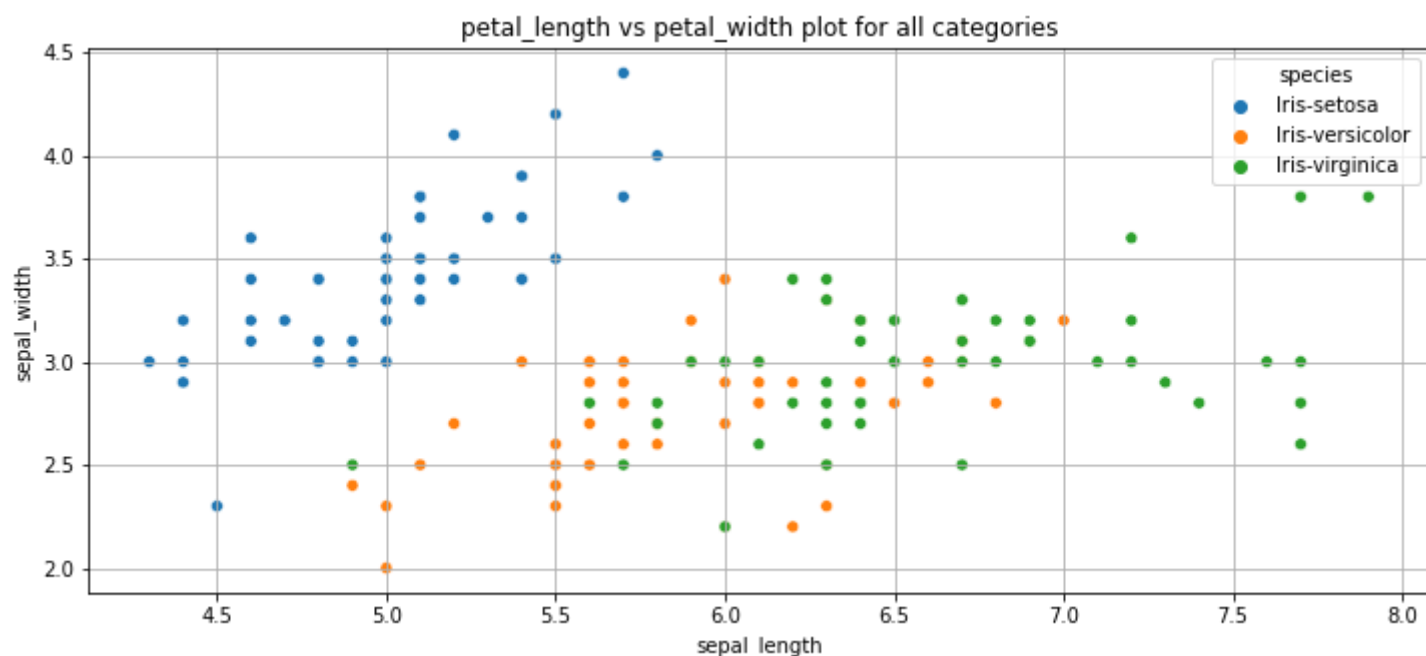
```
iris_tab=pd.crosstab(index=data['species'],columns='count')
```

```
iris_tab
```

col_0	count
species	
Iris-setosa	48
Iris-versicolor	50
Iris-virginica	49

## ▼ Step 3 : Data Visualization


```
#numeric v/s numeric v/s categorical
plt.figure(figsize=(12,5))
sns.scatterplot(x='sepal_length',y='sepal_width',data=data,hue='species')
plt.title('sepal_length vs sepal_width plot for all categories')
plt.grid()
plt.show()
```



Observations:

- Here the plot indicates that Iris-Setosa forms a cluster that can be separated from other categories based on sepal\_length and sepal\_width. While other 2 are mixed and cannot be separated directly

```
#numeric v/s numeric v/s categorical
plt.figure(figsize=(12,5))
sns.scatterplot(x='petal_length',y='petal_width',data=data,hue='species')
plt.title('petal_length vs petal_width plot for all categories')
plt.grid()
plt.show()
```




Observations:

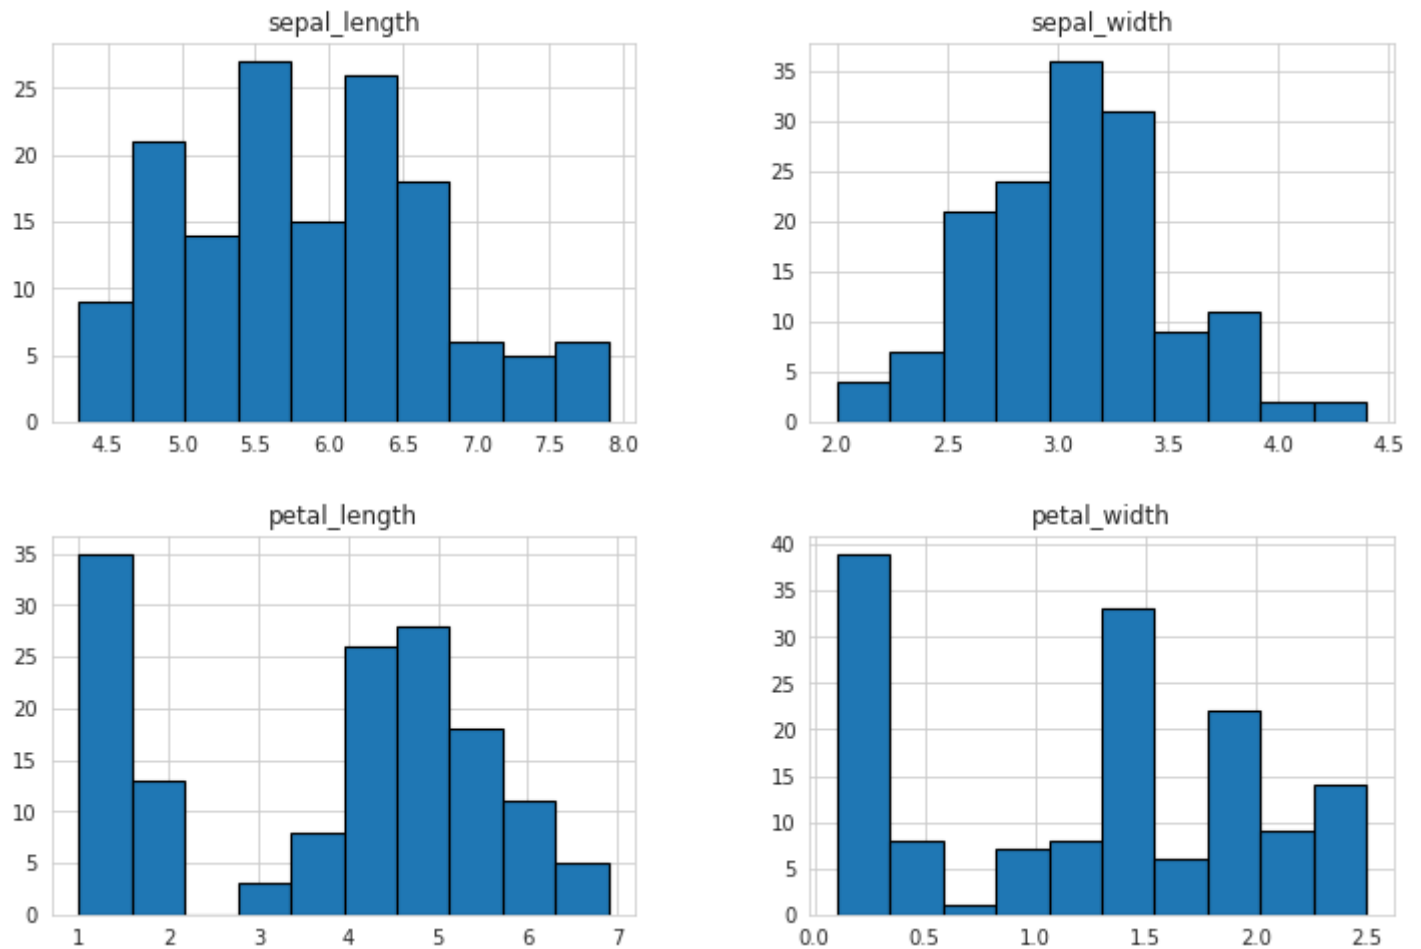
- Here also the plot indicates that Iris-Setosa forms a separate cluster that can be extracted from other categories based on petal\_length and petal\_width. While other 2 can also be separated .
- We can say that if the petal\_length < 2 and petal\_width < 0.95 the class is Setosa
- And if the petal\_length is between (2.9 and 5) and petal\_width is between (1 and 1.75) the class is Versicolor
- And if the petal\_length is between (4.75 and 7) and petal\_width is between (1.6 and 2.5) the class is Virginica

```
sns.set_style("whitegrid")
sns.pairplot(data,hue='species',size=3)
```

```
<seaborn.axisgrid.PairGrid at 0x7f077bfe4b70>
```



```
data.hist(edgecolor='black',linewidth=1.1,figsize=(12,8))
plt.show()
```




```
data.corr()
```

	sepal_length	sepal_width	petal_length	petal_width
sepal_length	1.000000	-0.109321	0.871305	0.817058
sepal_width	-0.109321	1.000000	-0.421057	-0.356376
petal_length	0.871305	-0.421057	1.000000	0.961883
petal_width	0.817058	-0.356376	0.961883	1.000000

```
sns.heatmap(data.corr(),annot=True,fmt='.2g',cmap='viridis')
```


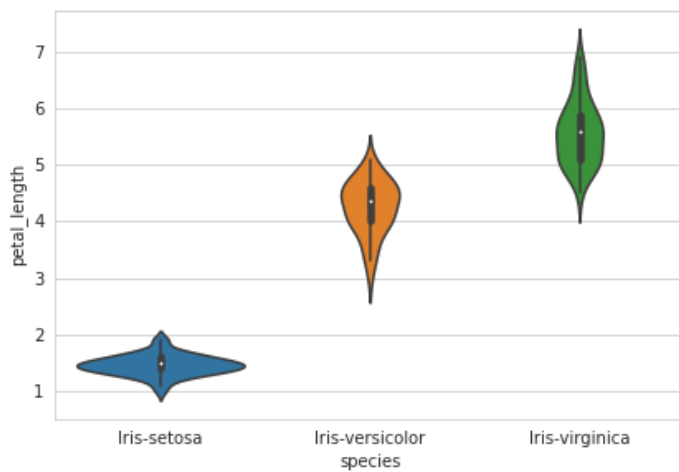
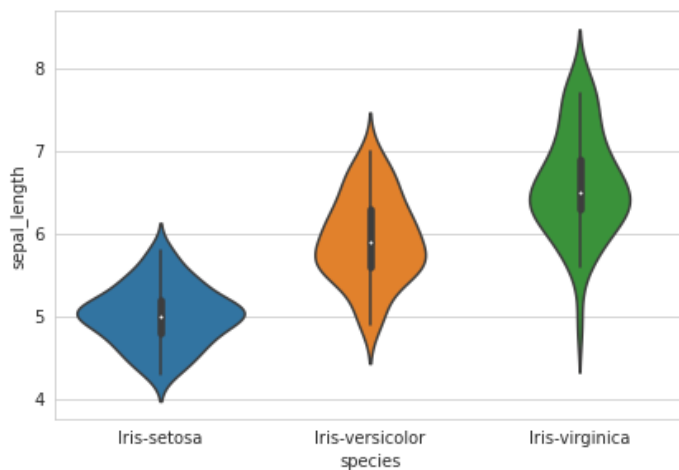
<matplotlib.axes.\_subplots.AxesSubplot at 0x7f07775d8898>



- petal\_length is highly correlated with sepal\_length and petal\_width and vice-versa.
- sepal\_width is not correlated with any feature it can be useful in separating features



```
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.violinplot(x='species',y='sepal_length',data=data)
plt.subplot(2,2,2)
sns.violinplot(x='species',y='sepal_width',data=data)
plt.subplot(2,2,3)
sns.violinplot(x='species',y='petal_length',data=data)
plt.subplot(2,2,4)
sns.violinplot(x='species',y='petal_width',data=data)
plt.show()
```



- The violinplot shows density of the length and width in the species. The thinner part denotes that

## ► K-Means Clustering

- K-Mean is a centroid-based algorithm, or a distance-based algorithm, where we calculate the distances to assign a point to a cluster. In K-Means, each cluster is associated with a centroid.


## How to Implementing K-Means Clustering ?

- Choose the number of clusters k
- Select k random points from the data as centroids
- Assign all the points to the closest cluster centroid
- Recompute the centroids of newly formed clusters
- Repeat steps 3 and 4

[ ] ↳ 2 cells hidden

## ▼ Using the elbow method to determine the optimal number of clusters for k-means clustering

```
plt.plot(range(1,15),loss)
plt.title('The elbow method')
plt.xlabel('Number of clusters')
plt.ylabel('Loss')
plt.show()
```



The bend in the curve is at 3.

```
kmeans=KMeans(n_clusters=3,random_state=95,n_jobs=-1)
kmeans.fit(x)
y_kmeans = kmeans.fit_predict(x)
```

y\_kmeans

```
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 2, 2, 2, 2, 0, 2, 2, 2,
       2, 2, 2, 0, 0, 2, 2, 2, 2, 0, 2, 0, 2, 0, 2, 2, 0, 0, 2, 2, 2, 2,
       2, 0, 2, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 2, 0], dtype=int32)
```

```
kmeans.cluster_centers_
```

```
array([[5.9016129 , 2.7483871 , 4.39354839, 1.43387097],
       [5.006      , 3.418      , 1.464      , 0.244      ],
       [6.85      , 3.07368421, 5.74210526, 2.07105263]])
```

```
#Visualising the clusters
```

```
plt.figure(figsize=(15,10))
```

```
plt.scatter(x[y_kmeans == 0, 0], x[y_kmeans == 0, 1], s = 100, c = 'purple', label = 'Iris-setosa')
```

```
plt.scatter(x[y_kmeans==1,0],x[y_kmeans==1,1],s=100,c='orange',label='Iris-versicolour')
```

```
plt.scatter(x[y_kmeans==2,0],x[y_kmeans==2,1],s=100,c='green',label='Iris-virginica')
```

```
#Plotting the centroids of the clusters
```

```
plt.scatter(kmeans.cluster_centers_[0],kmeans.cluster_centers_[1],s=100,c='red',label='Centroids')
```

```
plt.legend()
```

```
plt.show()
```



```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-97-2e192c0872cd> in <module>()
      1 #Visualising the clusters
      2 plt.figure(figsize=(15,10))
----> 3 plt.scatter(x[y_kmeans == 0, 0], x[y_kmeans == 0, 1], s = 100, c = 'purple', label =

```

## ▼ Prediction

```
data['predicted']=y_kmeans
```

```
/usr/local/lib/python3.6/dist-packages/pandas/core/indexes/base.py in get_loc(self, key,
```

```
data
```

	sepal_length	sepal_width	petal_length	petal_width	species	predicted
<b>0</b>	5.1	3.5	1.4	0.2	Iris-setosa	1
<b>1</b>	4.9	3.0	1.4	0.2	Iris-setosa	1
<b>2</b>	4.7	3.2	1.3	0.2	Iris-setosa	1
<b>3</b>	4.6	3.1	1.5	0.2	Iris-setosa	1
<b>4</b>	5.0	3.6	1.4	0.2	Iris-setosa	1
...	...	...	...	...	...	...
<b>145</b>	6.7	3.0	5.2	2.3	Iris-virginica	2
<b>146</b>	6.3	2.5	5.0	1.9	Iris-virginica	0
<b>147</b>	6.5	3.0	5.2	2.0	Iris-virginica	2
<b>148</b>	6.2	3.4	5.4	2.3	Iris-virginica	2
<b>149</b>	5.9	3.0	5.1	1.8	Iris-virginica	0

```
150 rows x 6 columns
```

```
False, False, False, True, False, False, False, True, False,
```

```
centroids=kmeans.cluster_centers_
```

```
centroids
```

```
array([[5.9016129 , 2.7483871 , 4.39354839, 1.43387097],
       [5.006      , 3.418      , 1.464      , 0.244      ],
       [6.85      , 3.07368421, 5.74210526, 2.07105263]])
```

```
df=data.copy()
```

```
df["species"]=df["species"].map({'Iris-versicolor':0,'Iris-setosa':1,'Iris-virginica':2}).astype(int)
```

```
df['predicted']=y_kmeans
```

```
df
```

	sepal_length	sepal_width	petal_length	petal_width	species	predicted
0	5.1	3.5	1.4	0.2	1	1
1	4.9	3.0	1.4	0.2	1	1
2	4.7	3.2	1.3	0.2	1	1
3	4.6	3.1	1.5	0.2	1	1
4	5.0	3.6	1.4	0.2	1	1
...	...	...	...	...	...	...

```
from sklearn.metrics import confusion_matrix
confusion_matrix(df['species'],y_kmeans)
```

```
array([[48,  0,  2],
       [ 0, 50,  0],
       [14,  0, 36]])
```

```
140      50      20      54      10      0      0
```


```
import scikitplot as skplt
true_labels = ["Iris-versicolor", "Iris-setosa", "Iris-virginica"]
pred_labels = ["Iris-versicolor", "Iris-setosa", "Iris-virginica"]
```

```
y_true=df['species']
```

```
data = confusion_matrix(y_true,y_kmeans)
df_cm = pd.DataFrame(data, columns=np.unique(y_true), index = np.unique(y_true))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (14,10))
sns.set(font_scale=1.4)
```

```
sns.heatmap(df_cm, cmap="Blues",xticklabels=true_labels,yticklabels=true_labels, annot=True,annot_kw
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f077763c828>
```



From the confusion matrix we can see that out of all 50 data of class Iris-versicolor 48 are correctly predicted whereas 2 are wrongly predicted as Iris-virginica. For 50 data of Iris-setosa all are correctly predicted. And for 50 data of Iris-virginica 36 are correctly predicted whereas 14 are wrongly predicted as Iris-versicolor.

```
data["predicted"]=data["predicted"].map({0:'Iris-versicolor',1:'Iris-setosa',2:'Iris-virginica'})
```

```
data.head()
```

	sepal_length	sepal_width	petal_length	petal_width	species	predicted
0	5.1	3.5	1.4	0.2	Iris-setosa	NaN
1	4.9	3.0	1.4	0.2	Iris-setosa	NaN
2	4.7	3.2	1.3	0.2	Iris-setosa	NaN
3	4.6	3.1	1.5	0.2	Iris-setosa	NaN
4	5.0	3.6	1.4	0.2	Iris-setosa	NaN

```
data.tail()
```

	sepal_length	sepal_width	petal_length	petal_width	species	predicted
145	6.7	3.0	5.2	2.3	Iris-virginica	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica	Iris-versicolor
147	6.5	3.0	5.2	2.0	Iris-virginica	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica	Iris-versicolor

```
#
```

Double-click (or enter) to edit

```
%%time
from sklearn.cluster import AgglomerativeClustering
from sklearn.metrics import silhouette_score
scores = []
for i in range(2,8):
    print(f'Starting Cluster #{i}')
    aggl_cluster=AgglomerativeClustering(n_clusters=i)
    aggl_cluster.fit(x)
    scores.append(silhouette_score(x,aggl_cluster.labels_,random_state=25))
```

```

Starting Cluster #2
Starting Cluster #3
Starting Cluster #4
Starting Cluster #5
Starting Cluster #6
Starting Cluster #7
CPU times: user 40.6 ms, sys: 27.1 ms, total: 67.7 ms
Wall time: 48.2 ms


```

```
# plot k vs Silhoutte Score
```

```

plt.plot(range(2,8), scores)
plt.xlabel('Clusters')
plt.ylabel('Silhoutte Score')
plt.title('Elbow Knee')
plt.legend('Knee')
plt.show()

```



```

aggl_cluster = AgglomerativeClustering(n_clusters=3)
aggl_cluster.fit(x)

```

```

AgglomerativeClustering(affinity='euclidean', compute_full_tree='auto',
connectivity=None, distance_threshold=None,
linkage='ward', memory=None, n_clusters=3)

```

```
aggl_cluster.labels_
```

```

array([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 2, 2, 2, 2, 0, 2, 2, 2,
        2, 2, 2, 0, 0, 2, 2, 2, 2, 0, 2, 0, 2, 0, 2, 2, 0, 0, 2, 2, 2, 2, 2,
        2, 0, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 0])

```

```
# Collect data points from each cluster
```

```
clusters_set = {i: np.where(aggl_cluster.labels_ == i)[0] for i in range(aggl_cluster.n_clusters)}
```

```
clusters_set
```

```
{0: array([ 50,  51,  52,  53,  54,  55,  56,  57,  58,  59,  60,  61,  62,
          63,  64,  65,  66,  67,  68,  69,  70,  71,  72,  73,  74,  75,
          76,  78,  79,  80,  81,  82,  83,  84,  85,  86,  87,  88,  89,
          90,  91,  92,  93,  94,  95,  96,  97,  98,  99, 101, 106, 113,
          114, 119, 121, 123, 126, 127, 133, 134, 138, 142, 146, 149]),
 1: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
          17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
          34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49]),
 2: array([ 77, 100, 102, 103, 104, 105, 107, 108, 109, 110, 111, 112, 115,
          116, 117, 118, 120, 122, 124, 125, 128, 129, 130, 131, 132, 135,
          136, 137, 139, 140, 141, 143, 144, 145, 147, 148])}]
```

```
agg_pred=aggl_cluster.fit_predict(x)
```

```
df1=data.copy()
df1["species"]=df1["species"].map({'Iris-versicolor':0,'Iris-setosa':1,'Iris-virginica':2}).astype(int)
df1['predicted']=agg_pred
df1
```

	sepal_length	sepal_width	petal_length	petal_width	species	predicted
<b>0</b>	5.1	3.5	1.4	0.2	1	1
<b>1</b>	4.9	3.0	1.4	0.2	1	1
<b>2</b>	4.7	3.2	1.3	0.2	1	1
<b>3</b>	4.6	3.1	1.5	0.2	1	1
<b>4</b>	5.0	3.6	1.4	0.2	1	1
...	...	...	...	...	...	...
<b>145</b>	6.7	3.0	5.2	2.3	2	2
<b>146</b>	6.3	2.5	5.0	1.9	2	0
<b>147</b>	6.5	3.0	5.2	2.0	2	2
<b>148</b>	6.2	3.4	5.4	2.3	2	2
<b>149</b>	5.9	3.0	5.1	1.8	2	0

150 rows × 6 columns


```
import scikitplot as skplt
true_labels = ["Iris-versicolor", "Iris-setosa", "Iris-virginica"]
pred_labels = ["Iris-versicolor", "Iris-setosa", "Iris-virginica"]

y_true=df1['species']

data = confusion_matrix(y_true,agg_pred)
df_cm = pd.DataFrame(data, columns=np.unique(y_true), index = np.unique(y_true))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (14,10))
sns.set(font_scale=1.4)

sns.heatmap(df_cm, cmap="Blues",xticklabels=true_labels,yticklabels=true_labels, annot=True,annot_kw
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f076f501668>



From the confusion matrix we can see that out of all 50 data of class Iris-versicolor 49 are correctly predicted whereas 2 are wrongly predicted as Iris-virginica. For 50 data of Iris-setosa all are wcorrectly predicted. And for 50 data of Iris-virginica 35 are correctly predicted whereas 15 are wrongly predicted as Iris-versicolor.

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.